

▼ Lab 1. PyTorch and ANNs

Deadline: Thursday, January 23, 11:59pm.

Total: 30 Points

Late Penalty: There is a penalty-free grace period of one hour past the deadline. Any work that is submitted between 1 hour and 24 hours past the deadline will receive a 20% grade deduction. No other late work is accepted. Quercus submission time will be used, not your local computer time. You can submit your labs as many times as you want before the deadline, so please submit often and early.

Grading TA: Kevin Course

This lab is partially based on an assignment developed by Prof. Jonathan Rose and Harris Chan.

This lab is a warm up to get you used to the PyTorch programming environment used in the course, and also to help you review and renew your knowledge of Python and relevant Python libraries. The lab must be done individually. Please recall that the University of Toronto plagiarism rules apply.

By the end of this lab, you should be able to:

1. Be able to perform basic PyTorch tensor operations.
2. Be able to load data into PyTorch
3. Be able to configure an Artificial Neural Network (ANN) using PyTorch
4. Be able to train ANNs using PyTorch
5. Be able to evaluate different ANN configurations

You will need to use numpy and PyTorch documentations for this assignment:

- <https://docs.scipy.org/doc/numpy/reference/>
- <https://pytorch.org/docs/stable/torch.html>

You can also reference Python API documentations freely.

What to submit

Submit a PDF file containing all your code, outputs, and write-up from parts 1-5. You can produce a PDF of your Google Colab file by going to **File > Print** and then save as PDF. The Colab instructions has more information.

Do not submit any other files produced by your code.

Include a link to your colab file in your submission.

Please use Google Colab to complete this assignment. If you want to use Jupyter Notebook, please complete the assignment and upload your Jupyter Notebook file to Google Colab for submission.

With Colab, you can export a PDF file using the menu option **File -> Print** and save as PDF file.

Colab Link

Submit make sure to include a link to your colab file here

Colab Link: <https://colab.research.google.com/drive/15AT4Qz4YkrflgdpM3gDd6gtK8mVWFgO9>

▼ Part 1. Python Basics [3 pt]

The purpose of this section is to get you used to the basics of Python, including working with functions, numbers, lists, and strings.

Note that we **will** be checking your code for clarity and efficiency.

If you have trouble with this part of the assignment, please review <http://cs231n.github.io/python-numpy-tutorial/>

▼ Part (a) -- 1pt

Write a function `sum_of_cubes` that computes the sum of cubes up to `n`. If the input to `sum_of_cubes` invalid (e.g. negative or non-integer `n`), the function should print out "Invalid input" and return `-1`.

```
1 def sum_of_cubes(n):  
2     """Return the sum of 1^3 + 2^3 + 3^3 + ... + n^3  
3     """
```

```

2     return the sum (1**3 + 2**3 + 3**3 + ... + n**3)
3
4     Precondition: n > 0, type(n) == int
5
6     >>> sum_of_cubes(3)
7     36
8     >>> sum_of_cubes(1)
9     1
10    """
11    if n < 0:
12        print("Invalid input")
13        return -1
14
15    sum = 0
16    for i in range(1, n + 1):
17        sum += i**3
18    return sum
19
20 #testing for correctness
21 print(sum_of_cubes(3))
22 print(sum_of_cubes(1))
23 print(sum_of_cubes(-10))

```

```

36
1
Invalid input
-1

```

▼ Part (b) -- 1pt

Write a function `word_lengths` that takes a sentence (string), computes the length of each word in that sentence, and returns the length of each word in a list. You can assume that words are always separated by a space character " " .

Hint: recall the `str.split` function in Python. If you aren't sure how this function works, try typing `help(str.split)` into a Python shell, or check out <https://docs.python.org/3.6/library/stdtypes.html#str.split>

```
1 help(str.split)
```

```
Help on method_descriptor:
```

```

split(...)
    S.split(sep=None, maxsplit=-1) -> list of strings

    Return a list of the words in S, using sep as the
    delimiter string.  If maxsplit is given, at most maxsplit
    splits are done.  If sep is not specified or is None, any
    whitespace string is a separator and empty strings are
    removed from the result.

```

```

1 def word_lengths(sentence):
2     """Return a list containing the length of each word in
3     sentence.
4
5     >>> word_lengths("welcome to APS360!")
6     [7, 2, 7]
7     >>> word_lengths("machine learning is so cool")
8     [7, 8, 2, 2, 4]
9     """
10    length_of_words = []
11
12    #in the case of an empty string, assume that there is only 1 word and it has length 0
13    if len(sentence) == 0:
14        length_of_words.append(0)
15        return length_of_words
16
17    words = sentence.split()
18
19    for word in words:
20        length_of_words.append(len(word))
21
22    return length_of_words
23
24 #testing for correctness

```

```

24 #testing for correctness
25 oneTest = word_lengths("welcome to APS360!")
26 print(oneTest)
27
28 anotherTest = word_lengths("machine learning is so cool")
29 print(anotherTest)
30
31 oneMoreTest = word_lengths("")
32 print(oneMoreTest)
33

```

```

[7, 2, 7]
[7, 8, 2, 2, 4]
[0]

```

▼ Part (c) – 1pt

Write a function `all_same_length` that takes a sentence (string), and checks whether every word in the string is the same length. You should call the function `word_lengths` in the body of this new function.

```

1 def all_same_length(sentence):
2     """Return True if every word in sentence has the same
3     length, and False otherwise.
4
5     >>> all_same_length("all same length")
6     False
7     >>> word_lengths("hello world")
8     True
9     """
10    #in the case of an empty string, assume that all words have the same length and return True
11
12    length_of_words = word_lengths(sentence)
13    first_length = length_of_words[0]
14
15    for length in length_of_words:
16        if length != first_length:
17            return False
18    return True
19
20 #function to print out value of bool
21 def print_bool(TF):
22     if TF:
23         print("True")
24     else:
25         print("False")
26
27 #testing for correctness
28 firstTest = all_same_length("all same length")
29 print_bool(firstTest)
30
31 secondTest = all_same_length("hello world")
32 print_bool(secondTest)
33
34 thirdTest = all_same_length("")
35 print_bool(thirdTest)

```

```

False
True
True

```

▼ Part 2. NumPy Exercises [5 pt]

In this part of the assignment, you'll be manipulating arrays using NumPy. Normally, we use the shorter name `np` to represent the package `numpy`.

```

1 import numpy as np

```

▼ Part (a) – 1pt

The below variables `matrix` and `vector` are numpy arrays. Explain what you think `<NumpyArray>.size` and `<NumpyArray>.shape` represent.

```
1 matrix = np.array([[1., 2., 3., 0.5],
2                   [4., 5., 0., 0.],
3                   [-1., -2., 1., 1.]])
4 vector = np.array([2., 0., 1., -2.])
```

```
1 matrix.size
2 # determines the number of elements in the array
```

```
↳ 12
```

```
1 matrix.shape
2 # determines the number of rows and columns, output is represented as a tuple (number of rows, number of columns)
```

```
↳ (3, 4)
```

```
1 vector.size
2 # determines the number of elements in the vector
```

```
↳ 4
```

```
1 vector.shape
2 # determines the number of rows in a row vector or the number of columns in a column vector
3 # but vectors are one-dimensional in NumPy
```

```
↳ (4,)
```

▼ Part (b) -- 1pt

Perform matrix multiplication `output = matrix x vector` by using for loops to iterate through the columns and rows. Do not use any builtin NumPy functions. Cast your output into a NumPy array, if it isn't one already.

Hint: be mindful of the dimension of output

```
1 matrix_dim = matrix.shape #(3,4)
2 vector_dim = vector.shape #(4,)
3
4
5 output = np.array([0., 0., 0.]) # 3x4 array * 4x1 vector = 3x1 array
6 sum = 0
7
8 if vector_dim[0] == matrix_dim[1]:
9     for row in range(0, matrix_dim[0]):
10         for col in range(0, matrix_dim[1]):
11             output[row] += vector[col] * matrix[row][col]
12
```

```
1 output
```

```
↳ array([ 4.,  8., -3.])
```

▼ Part (c) -- 1pt

Perform matrix multiplication `output2 = matrix x vector` by using the function `numpy.dot`.

We will never actually write code as in part(c), not only because `numpy.dot` is more concise and easier to read/write, but also performance-wise `numpy.dot` is much faster (it is written in C and highly optimized). In general, we will avoid for loops in our code.

```
1 output2 = np.dot(matrix, vector)
```

```
1 output2
```

```
↳ array([ 4.,  8., -3.])
```

▼ Part (d) -- 1pt

As a way to test for consistency, show that the two outputs match.

```

1 if np.array_equal(output, output2):
2     print("Outputs match")
3 else:
4     print("Outputs don't match")

```

☞ Outputs match

▼ Part (e) -- 1pt

Show that using `np.dot` is faster than using your code from part (c).

You may find the below code snippet helpful:

```

1 import time
2 start_time_my_function = time.time()
3
4 matrix_dim = matrix.shape #(3,4)
5 vector_dim = vector.shape #(4,)
6
7 output = np.array([0., 0., 0.])
8 sum = 0
9
10 if vector_dim[0] == matrix_dim[1]:
11     for row in range(0, matrix_dim[0]):
12         for col in range(0, matrix_dim[1]):
13             output[row] += vector[col] * matrix[row][col]
14
15 end_time_my_function = time.time()
16
17 print("Runtime of function in part (b): {}".format(end_time_my_function - start_time_my_function))
18
19 start_time_numpy = time.time()
20 outputX = np.dot(matrix, vector)
21 end_time_numpy = time.time()
22
23 print("Runtime of NumPy function .dot: {}".format(end_time_numpy - start_time_numpy))
24
25 # compute the difference
26 #diff = abs((end_time_my_function - start_time_my_function) - (end_time_numpy - start_time_numpy))
27 #print("np.dot is {} seconds faster than the function in part (b)".format(diff))

```

☞ Runtime of function in part (b): 0.0002918243408203125
Runtime of NumPy function .dot: 8.58306884765625e-05

▼ Part 3. Images [6 pt]

A picture or image can be represented as a NumPy array of “pixels”, with dimensions $H \times W \times C$, where H is the height of the image, W is the width of the image, and C is the number of colour channels. Typically we will use an image with channels that give the the Red, Green, and Blue “level” of each pixel, which is referred to with the short form RGB.

You will write Python code to load an image, and perform several array manipulations to the image and visualize their effects.

```

1 import matplotlib.pyplot as plt

```

▼ Part (a) -- 1 pt

This is a photograph of a dog whose name is Mochi.



Load the image from its url (https://drive.google.com/uc?export=view&id=1oaLVR2hr1_qzpKQ47i9rVUIklwbDcews) into the variable `img` using the `plt.imread` function.

Hint: You can enter the URL directly into the `plt.imread` function as a Python string.

```
1 img = plt.imread("https://drive.google.com/uc?export=view&id=1oaLVR2hr1_qzpKQ47i9rVUIklwbDcews")
```


▼ Part (b) -- 1pt

Use the function `plt.imshow` to visualize `img`.

This function will also show the coordinate system used to identify pixels. The origin is at the top left corner, and the first dimension indicates the Y (row) direction, and the second dimension indicates the X (column) dimension.

```
1 plt.imshow(img)
```

```
<matplotlib.image.AxesImage at 0x7fc971833fd0>
```




▼ Part (c) -- 2pt

Modify the image by adding a constant value of 0.25 to each pixel in the `img` and store the result in the variable `img_add`. Note that, since the range for the pixels needs to be between [0, 1], you will also need to clip `img_add` to be in the range [0, 1] using `numpy.clip`. Clipping sets any value that is outside of the desired range to the closest endpoint. Display the image using `plt.imshow`.

```
1 img_add = np.clip(img + 0.25, 0, 1)
2 plt.imshow(img_add)
```

```
<matplotlib.image.AxesImage at 0x7fc971833fd0>
```

<matplotlib.image.AxesImage at 0x7fc96ff469b0>




▼ Part (d) -- 2pt

Crop the **original** image (`img` variable) to a 130 x 150 image including Mochi's face. Discard the alpha colour channel (i.e. resulting `img_cropped` should **only have RGB channels**)

Display the image.

```
1 img_cropped = img[0:130, 0:150, 0:3]
2 plt.imshow(img_cropped)
```

<matplotlib.image.AxesImage at 0x7fc96ff2c5c0>



▼ Part 4. Basics of PyTorch [6 pt]

PyTorch is a Python-based neural networks package. Along with tensorflow, PyTorch is currently one of the most popular machine learning libraries.

PyTorch, at its core, is similar to Numpy in a sense that they both try to make it easier to write codes for scientific computing achieve improved performance over vanilla Python by leveraging highly optimized C back-end. However, compare to Numpy, PyTorch offers much better GPU support and provides many high-level features for machine learning. Technically, Numpy can be used to perform almost every thing PyTorch does. However, Numpy would be a lot slower than PyTorch, especially with CUDA GPU, and it would take more effort to write machine learning related code compared to using PyTorch.

```
1 import torch
```

▼ Part (a) -- 1 pt

Use the function `torch.from_numpy` to convert the numpy array `img_cropped` into a PyTorch tensor. Save the result in a variable called `img_torch`.

```
1 img_torch = torch.from_numpy(img_cropped)
```

▼ Part (b) -- 1pt

Use the method `<Tensor>.shape` to find the shape (dimension and size) of `img_torch`.

```
1 img_torch.shape
```

```
↳ torch.Size([130, 150, 3])
```

▼ Part (c) – 1pt

How many floating-point numbers are stored in the tensor `img_torch`?

```
1 img_torch.shape[0]*img_torch.shape[1]*img_torch.shape[2]
```

```
↳ 58500
```

▼ Part (d) – 1 pt

What does the code `img_torch.transpose(0,2)` do? What does the expression return? Is the original variable `img_torch` updated? Explain.

```
1 #img_torch.transpose(0,2) transposes the rows and columns of img_torch by swapping the dimensions 0 and 2
2 #it returns a tensor that is the transpose of img_torch
3 #img_torch is not updated
4
5 print("dimensions of img_torch before transpose: {}".format(img_torch.shape))
6 print("dimensions of operation transposed img_torch: {}".format(img_torch.transpose(0,2).shape))
7 print("dimensions of img_torch after transpose: {}".format(img_torch.shape))
```

```
↳ dimensions of img_torch before transpose: torch.Size([130, 150, 3])
   dimensions of operation transposed img_torch: torch.Size([3, 150, 130])
   dimensions of img_torch after transpose: torch.Size([130, 150, 3])
```

▼ Part (e) – 1 pt

What does the code `img_torch.unsqueeze(0)` do? What does the expression return? Is the original variable `img_torch` updated? Explain.

```
1 #img_torch.unsqueeze(0) will return a new tensor with a dimension inserted where specified in the function call
2 #it returns a new tensor with an additional dimension
3 #img_torch is not updated
4
5 print("dimensions of img_torch before unsqueeze: {}".format(img_torch.shape))
6 print("dimensions of operation img_torch.unsqueeze: {}".format(img_torch.unsqueeze(0).shape))
7 print("dimensions of img_torch after unsqueeze: {}".format(img_torch.shape))
8
```

```
↳ dimensions of img_torch before unsqueeze: torch.Size([130, 150, 3])
   dimensions of operation img_torch.unsqueeze: torch.Size([1, 130, 150, 3])
   dimensions of img_torch after unsqueeze: torch.Size([130, 150, 3])
```

▼ Part (f) – 1 pt

Find the maximum value of `img_torch` along each colour channel? Your output should be a one-dimensional PyTorch tensor with exactly three values.

Hint: lookup the function `torch.max`.

```
1 torch.tensor([torch.max(img_torch[:, :, 0]), torch.max(img_torch[:, :, 1]), torch.max(img_torch[:, :, 2])])
```

```
↳ tensor([0.8941, 0.7882, 0.6745])
```

▼ Part 5. Training an ANN [10 pt]

The sample code provided below is a 2-layer ANN trained on the MNIST dataset to identify digits less than 3 or greater than and equal to 3. Modify the code by changing any of the following and observe how the accuracy and error are affected:

- number of training iterations
- number of hidden units
- numbers of layers
- types of activation functions
- learning rate


```

1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4 from torchvision import datasets, transforms
5 import matplotlib.pyplot as plt # for plotting
6 import torch.optim as optim
7
8 torch.manual_seed(1) # set the random seed
9
10 # define a 2-layer artificial neural network
11 class Pigeon(nn.Module):
12     def __init__(self):
13         super(Pigeon, self).__init__()
14         self.layer1 = nn.Linear(28 * 28, 30) # 28*28 node -> 30 node
15
16         # change the number of hidden units
17         #self.layer1 = nn.Linear(28 * 28, 3)
18         #self.layer2 = nn.Linear(3, 1)
19
20         # add another layer
21         #self.layer3 = nn.Linear(30,30) #intermediate 30 node input -> 30 node output layer
22         #self.layer4 = nn.Linear(30,30)
23         #self.layer5 = nn.Linear(30,30)
24         #self.layer6 = nn.Linear(30,30)
25         #self.layer7 = nn.Linear(30,30)
26         #self.layer8 = nn.Linear(30,30)
27
28         self.layer2 = nn.Linear(30, 1) # 30 node -> 1 node
29
30     def forward(self, img):
31         flattened = img.view(-1, 28 * 28)
32         activation1 = self.layer1(flattened)
33         activation1 = F.relu(activation1)
34
35         #change the activation function
36         #activation1 = F.sigmoid(activation1)
37         #activation1 = F.tanh(activation1)
38
39         # add another layer
40         #activation3 = self.layer3(activation1) # after adding intermediate layer between input and hidden layer
41         #activation4 = self.layer4(activation1)
42         #activation5 = self.layer5(activation1)
43         #activation6 = self.layer6(activation1)
44         #activation7 = self.layer7(activation1)
45         #activation8 = self.layer8(activation1)
46
47         activation2 = self.layer2(activation1)
48         return activation2
49
50 pigeon = Pigeon()
51
52 # load the data
53 mnist_data = datasets.MNIST('data', train=True, download=True)
54 mnist_data = list(mnist_data)
55 mnist_train = mnist_data[:1000]
56 mnist_val = mnist_data[1000:2000]
57 img_to_tensor = transforms.ToTensor()
58
59
60 # simplified training code to train `pigeon` on the "small digit recognition" task
61 criterion = nn.BCEWithLogitsLoss()
62 optimizer = optim.SGD(pigeon.parameters(), lr=0.005, momentum=0.9)
63
64 # change learning rate
65 #optimizer = optim.SGD(pigeon.parameters(), lr=0.01, momentum=0.9)
66
67 # change the number of training iterations
68 #iterations = 2
69
70 #for iter in range(iterations) :
71 for (image, label) in mnist_train:
72     # actual ground truth: is the digit less than 3?
73     actual = torch.tensor(label < 3).reshape([1, 1]).type(torch.FloatTensor)

```

```

73 actual = criterion(out, actual).type(torch.FloatTensor)
74 # pigeon prediction
75 out = pigeon(img_to_tensor(image)) # step 1-2
76 # update the parameters based on the loss
77 loss = criterion(out, actual) # step 3
78 loss.backward() # step 4 (compute the updates for each parameter)
79 optimizer.step() # step 4 (make the updates for each parameter)
80 optimizer.zero_grad() # a clean up step for PyTorch
81
82 # computing the error and accuracy on the training set
83 error = 0
84 for (image, label) in mnist_train:
85     prob = torch.sigmoid(pigeon(img_to_tensor(image)))
86     if (prob < 0.5 and label < 3) or (prob >= 0.5 and label >= 3):
87         error += 1
88 print("Training Error Rate:", error/len(mnist_train))
89 print("Training Accuracy:", 1 - error/len(mnist_train))
90
91
92 # computing the error and accuracy on a test set
93 error = 0
94 for (image, label) in mnist_val:
95     prob = torch.sigmoid(pigeon(img_to_tensor(image)))
96     if (prob < 0.5 and label < 3) or (prob >= 0.5 and label >= 3):
97         error += 1
98 print("Test Error Rate:", error/len(mnist_val))
99 print("Test Accuracy:", 1 - error/len(mnist_val))

```

```

❏ Training Error Rate: 0.036
Training Accuracy: 0.964
Test Error Rate: 0.079
Test Accuracy: 0.921

```

▼ Part (a) – 3 pt

Comment on which of the above changes resulted in the best accuracy on training data? What accuracy were you able to achieve?

```

1 """
2 A list of hyperparameters changed and the resulting training and testing accuracies is included in the Appendix after part (c)
3
4 The highest accuracy in the training data occurred in the test cases when:
5 - the number of training iterations increased from 1 to 50
6   - Training Accuracy: 1.0
7 - the number of hidden units increased from 30 to 60 and 300
8   - Training Accuracy: 0.969 and Training Accuracy: 0.977, respectively
9 - the number of layers increased from 2 to 5
10  - Training Accuracy: 0.965
11
12 Increasing the above three hyperparameters allows the NN to become more familiar with the training dataset because it has more
13 passes over the data, allowing more opportunities to find patterns and other trends in the data. However, allowing many passes
14 over the training data will also cause the NN to overfit to the noise and other irregularities present in the training dataset,
15 which may not be present in the testing dataset or other unseen datasets. Overfitting may be why the Training Accuracy decreased
16 when the number of layers used increased from 5 to 8 (0.965 -> 0.954)
17 """

```

▼ Part (b) – 3 pt

Comment on which of the above changes resulted in the best accuracy on testing data? What accuracy were you able to achieve?

```

1 """
2 The highest accuracy in the testing data occurred in the test cases when:
3 - the number of training iterations is 2
4   - Test Accuracy: 0.943
5 - the number of hidden units is 300
6   - Test Accuracy: 0.929
7 - the number of layers is 5
8   - Test Accuracy: 0.9299999999999999
9
10 If increasing a NN hyperparameter increased the accuracy of the NN on the training data, then it is reasonable to expect that
11 such changes would also result in improved accuracy on the testing data. This is the case when the number of hidden units is
12 set to 300 and the number of layers is set to 5.
13

```

```

14 However, it is interesting to note that the test accuracy is greatest when the number of training iterations is set to 2, not
15 when it is set to 50 as in part (a). This implies that the NN experienced overfitting when the number of iterations was 50, and
16 it began adapting to the noise that was present in the training dataset, and trying to fit this noise onto the testing dataset.
17 Because the testing dataset lacked the noise present in the training dataset, it was difficult for the NN to apply this pattern
18 recognition to something that doesn't exist.
19 """

```

▼ Part (c) – 4 pt

Which model hyperparameters should you use, the ones from (a) or (b)?

```

1 """
2 To achieve the highest accuracy on an unseen data set, we should:
3 - set the number of training iterations to 2 (part (b))
4 - set the number of hidden units to 300 (part (a))
5 - set the number of layers to 5 (part (a))
6 - keep all other hyperparameters the same as in the original model
7
8 In choosing the values of these hyperparameters, it is important to balance too few passes over the data against too many passes.
9 With too few passes, the NN does not have an opportunity to detect patterns, learn of their frequencies, and predict their
10 occurrence in future samples. However with too many passes, the NN will become overfitted to the training data, and it will
11 incorporate irregularities and noise present in the training data into its predictions on future samples where the same noise
12 may not exist.
13
14 It is also important to balance the accuracy of the NN and its performance. Although the test accuracy is higher when the number
15 of hidden units is set to 300, the accuracy only improved by 0.008 when compared to the accuracy of the original model that used
16 30 hidden units. For an accuracy improvement of less than 1%, it can be argued that the slight improvement in accuracy is not
17 a good justification for increasing the runtime.
18 """

```

▼ Appendix: Changing hyperparameters for Part 5

```

1 """
2 I changed one parameter at a time
3 *****
4         Original
5 number of training iterations: 1
6 number of hidden units: 30
7 numbers of layers: 2
8 types of activation functions: ReLU
9 learning rate: 0.005
10 number of training points: 1000
11 number of validation points: 1000
12
13 Training Error Rate: 0.036
14 Training Accuracy: 0.964
15 Test Error Rate: 0.079
16 Test Accuracy: 0.921
17
18 *****
19         Changed number of training iterations
20 number of training iterations: 2
21
22 Training Error Rate: 0.016
23 Training Accuracy: 0.984
24 Test Error Rate: 0.057
25 Test Accuracy: 0.943
26
27 number of training iterations: 10
28
29 Training Error Rate: 0.001
30 Training Accuracy: 0.999
31 Test Error Rate: 0.058
32 Test Accuracy: 0.942
33
34 number of training iterations: 50
35
36 Training Error Rate: 0.0
37 Training Accuracy: 1.0
38 Test Error Rate: 0.059

```

```
39 Test Accuracy: 0.9410000000000001
40
41 *****
42         Changed number of hidden units
43 number of hidden units: 3
44
45 Training Error Rate: 0.103
46 Training Accuracy: 0.897
47 Test Error Rate: 0.143
48 Test Accuracy: 0.857
49
50 number of hidden units: 60
51
52 Training Error Rate: 0.031
53 Training Accuracy: 0.969
54 Test Error Rate: 0.08
55 Test Accuracy: 0.92
56
57 number of hidden units: 300
58
59 Training Error Rate: 0.023
60 Training Accuracy: 0.977
61 Test Error Rate: 0.071
62 Test Accuracy: 0.929
63
64 *****
65         Changed number of layers
66 numbers of layers: 3
67
68 Training Error Rate: 0.043
69 Training Accuracy: 0.957
70 Test Error Rate: 0.081
71 Test Accuracy: 0.919
72
73 numbers of layers: 5
74
75 Training Error Rate: 0.035
76 Training Accuracy: 0.965
77 Test Error Rate: 0.07
78 Test Accuracy: 0.9299999999999999
79
80 numbers of layers: 8
81
82 Training Error Rate: 0.046
83 Training Accuracy: 0.954
84 Test Error Rate: 0.086
85 Test Accuracy: 0.914
86
87 *****
88         Changed activation function
89 type of activation function: sigmoid
90
91 Training Error Rate: 0.073
92 Training Accuracy: 0.927
93 Test Error Rate: 0.117
94 Test Accuracy: 0.883
95
96 type of activation function: tanh
97
98 Training Error Rate: 0.04
99 Training Accuracy: 0.96
100 Test Error Rate: 0.094
101 Test Accuracy: 0.906
102
103 type of activation function: softmax
104
105 Training Error Rate: 0.312
106 Training Accuracy: 0.688
107 Test Error Rate: 0.297
108 Test Accuracy: 0.7030000000000001
109
110 type of activation function: softsign
111
112 Training Error Rate: 0.053
```

```
112 Training Error Rate: 0.055
113 Training Accuracy: 0.947
114 Test Error Rate: 0.108
115 Test Accuracy: 0.892
116
117 *****
118         Changed learning rate
119 learning rate: 0.00001
120
121 Training Error Rate: 0.321
122 Training Accuracy: 0.679
123 Test Error Rate: 0.311
124 Test Accuracy: 0.6890000000000001
125
126 learning rate: 0.0001
127
128 Training Error Rate: 0.312
129 Training Accuracy: 0.688
130 Test Error Rate: 0.297
131 Test Accuracy: 0.7030000000000001
132
133 learning rate: 0.01
134
135 Training Error Rate: 0.039
136 Training Accuracy: 0.961
137 Test Error Rate: 0.082
138 Test Accuracy: 0.918
139
140 learning rate: 0.1
141
142 Training Error Rate: 0.312
143 Training Accuracy: 0.688
144 Test Error Rate: 0.297
145 Test Accuracy: 0.7030000000000001
146
147 *****
148         Changed number of points in training and validation datasets
149 number of training points: 500
150 number of validation points: 1500
151
152 Training Error Rate: 0.06
153 Training Accuracy: 0.94
154 Test Error Rate: 0.11266666666666666
155 Test Accuracy: 0.8873333333333333
156
157 number of training points: 1500
158 number of validation points: 500
159
160 Training Error Rate: 0.05
161 Training Accuracy: 0.95
162 Test Error Rate: 0.074
163 Test Accuracy: 0.926
164
165 number of training points: 1800
166 number of validation points: 200
167
168 Training Error Rate: 0.055
169 Training Accuracy: 0.945
170 Test Error Rate: 0.085
171 Test Accuracy: 0.915
172 ""
```